2018

Global Azure BOOTCAMP

VERONA

Gli sponsor Mondiali

Microsoft

Gli sponsor del nostro evento

Platinum Sponsor

Gli sponsor del nostro evento

Gold Sponsor

Gli sponsor del nostro evento

digital joy

dgroove
Tecnologia emozionale.

Basic Sponsor

Tweet della giornata

#GlobalAzure

@cloudgen_verona

2018 Global Azure BOOTCAMP

ARGOMENTO

Azure Cosmos DB

Un moderno database nel cloud

MARTINO BORDIN

@martinobordin

martinobordin

martinobordin

What is Azure Cosmos DB

https://azure.microsoft.com/en-us/services/cosmos-db/

«Globally distributed, multi-model database service»

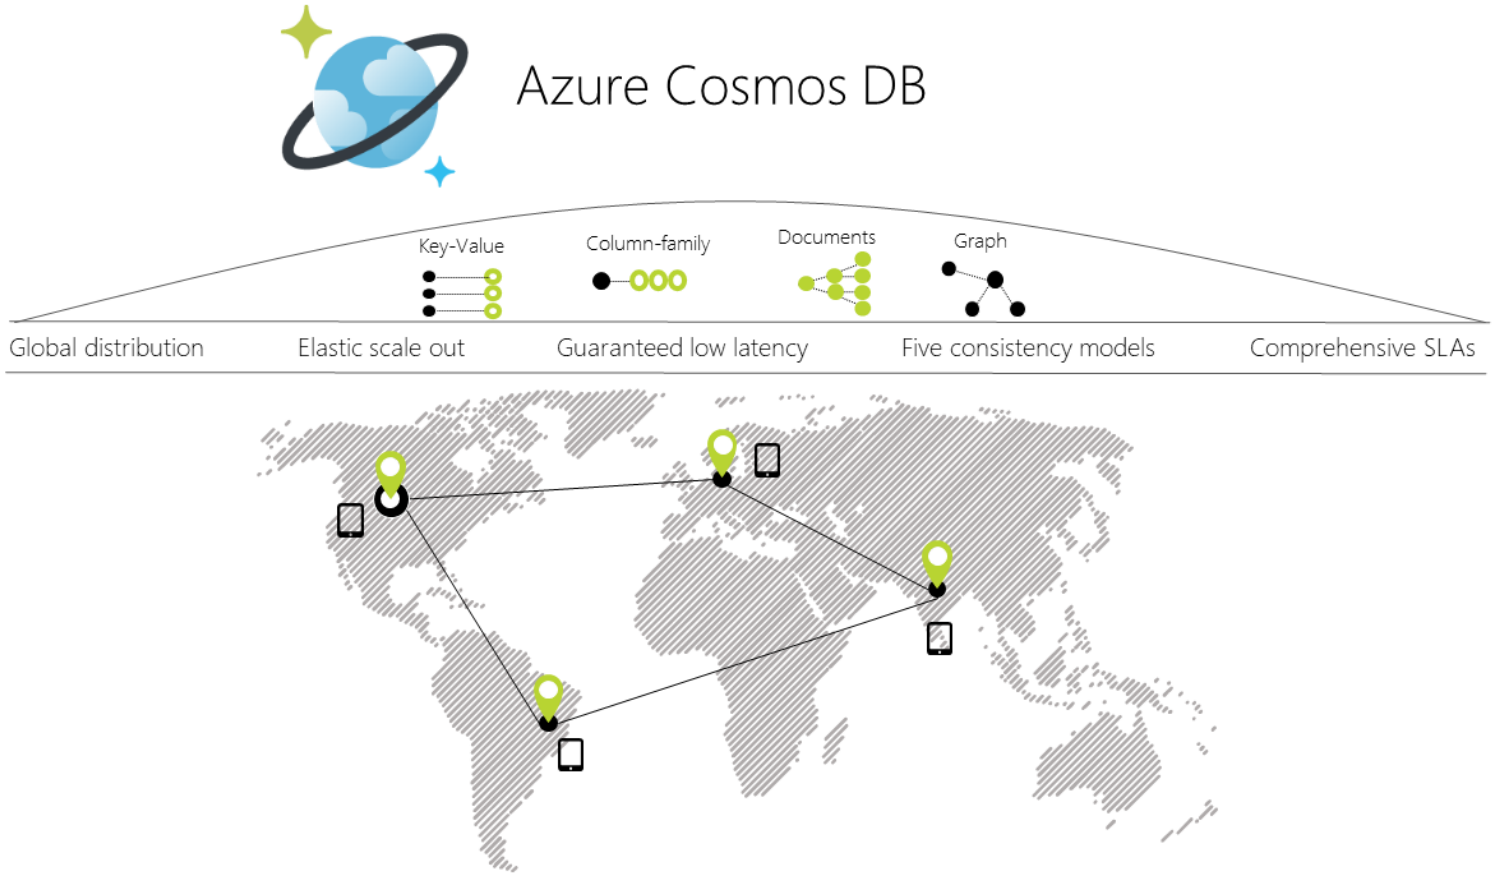

To SQL or to NoSQL?

- Schemaless
- Horizontal scaling
- Consistency model -> ACID vs Eventual Consistency
- Design & Development
- Alternatives:
  - Key-value: *Redis*
  - Wide-column: *Cassandra*
  - Document: *MongoDb, RavenDb*
  - Graph: *Neo4J*

DocumentDB features

Document database on Cosmos DB

- Collections, documents..and more!
- SLA, Security & Logging
- Request Unit & Costs
- Partitions
- Different consistency levels
- Customizable indexing
- Using SQL with NoSQL
- TTL
- Global distribution
- Development
- Emulator
- Tools
- Backup/Restore
- What's next

Collections, documents..and more!

- Databases
- Collections (Tables)
  - Stored Procedures
  - Triggers
  - User defined functions
- Documents (*Rows/Columns*)
- Attachments

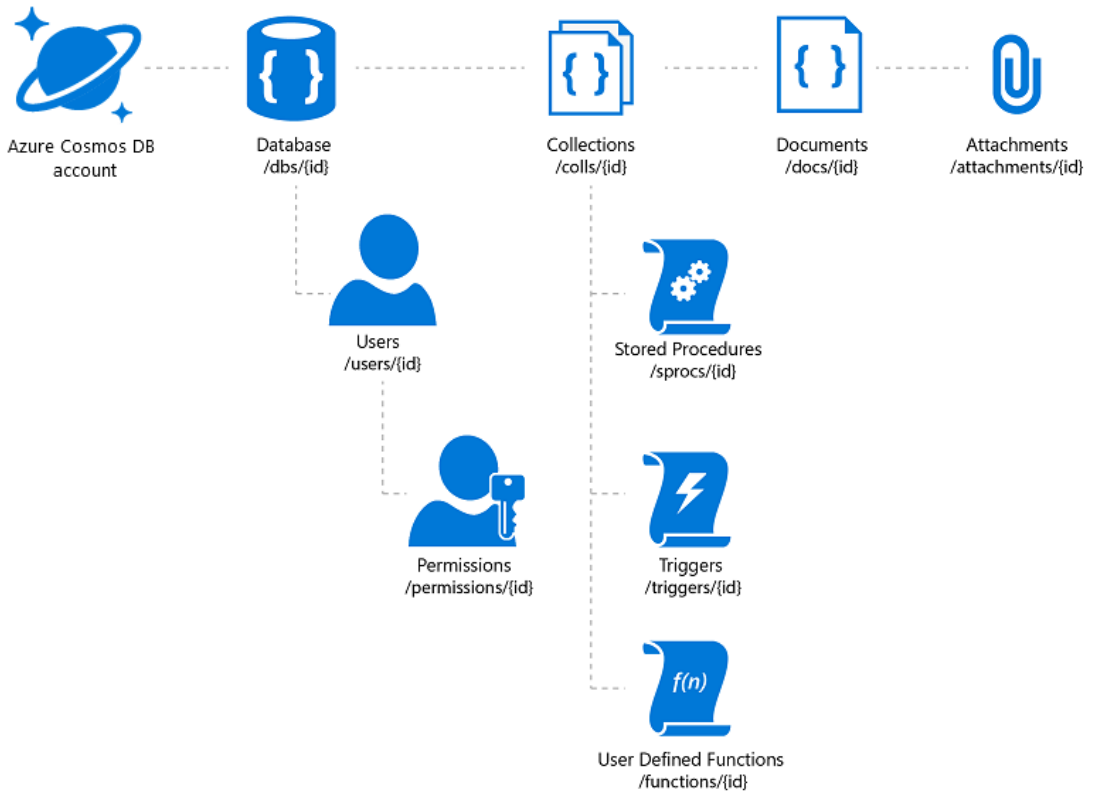

SLA, Security & Logging

AVAILABILITY (% MONTH)	SERVICE CREDIT
< 99.99%	10%
< 99%	25%

- SLA
  - 99.999% availability
  - <10 ms reads <15 ms writes
- Security
  - Key authentication:
    - *Master keys* -> Administrative resources (Databases, Users, Permissions) -> **GLOBAL**
    - *Resource tokens* -> Application resources : (Collections, Documents, Attachments, SP, TRI, UDF) -> **USER**
  - **IP** Access Control Policy
  - Data encrypted **in transit** and **at rest**
- Logging
  - CRUD, Database, Access
  - Storage Account, Event HUB, Log Analytics

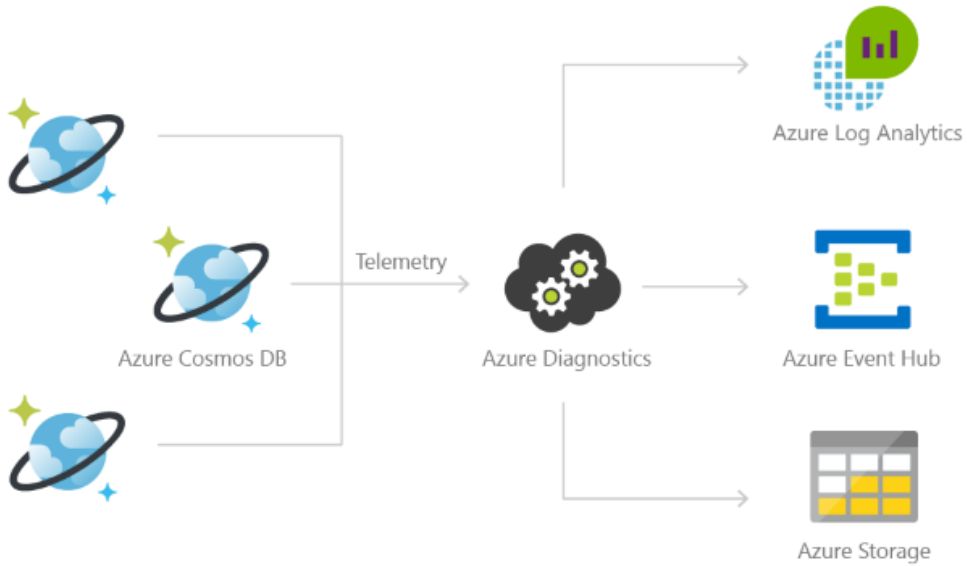

Request Unit & Costs

- Request Unit (https://www.documentdb.com/capacityplanner)
  - Performance is based on **provisioned Request Units/sec**: 1 KB Doc -> 1 RU Read / 5 RU Write
  - Capacity planner

- Costs (https://azure.microsoft.com/en-us/pricing/calculator)
  - SSD Storage -> $0.25 GB/month
  - 100 RUs, (min **400 RUs) ->**$0.008/hour

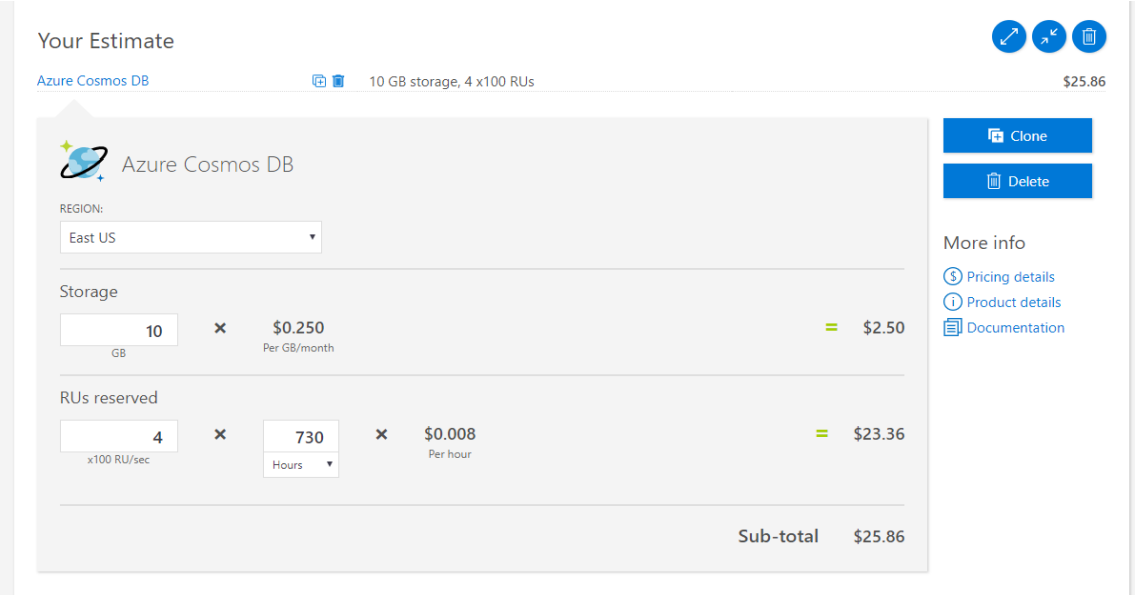

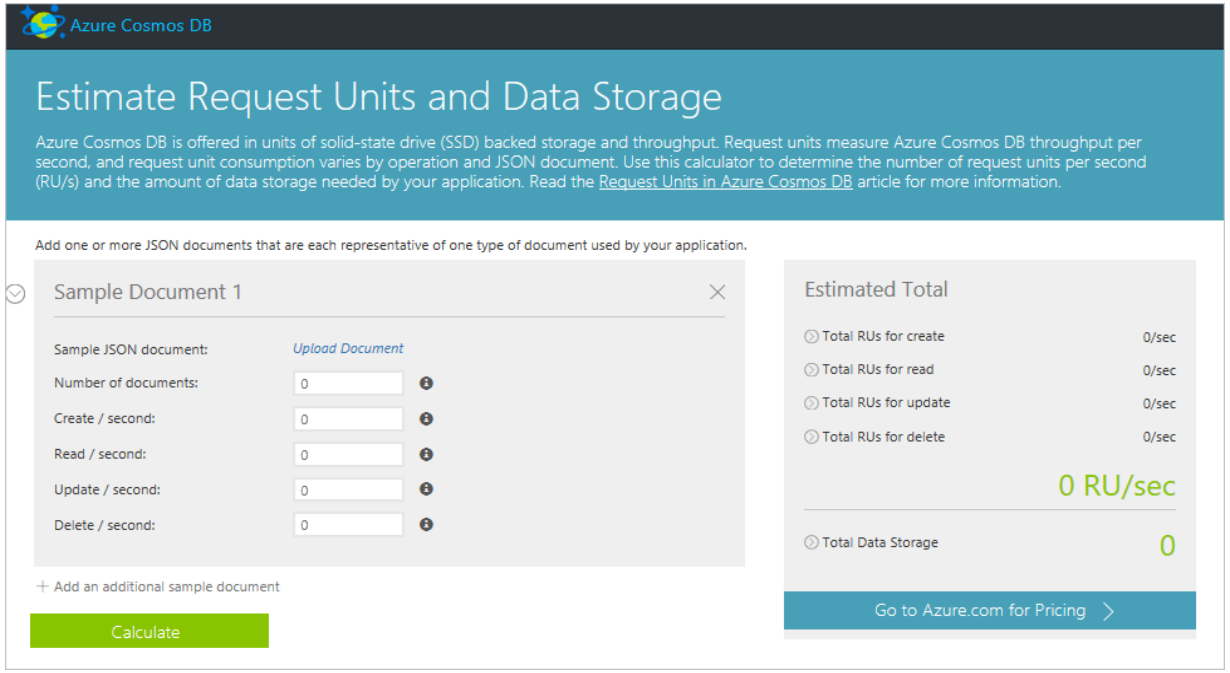

Partitions

2 container (collections) types:

- Fixed:
  - Min 400 RU/S
  - Max 10,000 RU/s
  - Max 10 GB
- Unlimited
  - Min 1,000 RU/s
  - Max (?) 1,000,000 RU/s
  - **Must have a <u>partition key</u>**

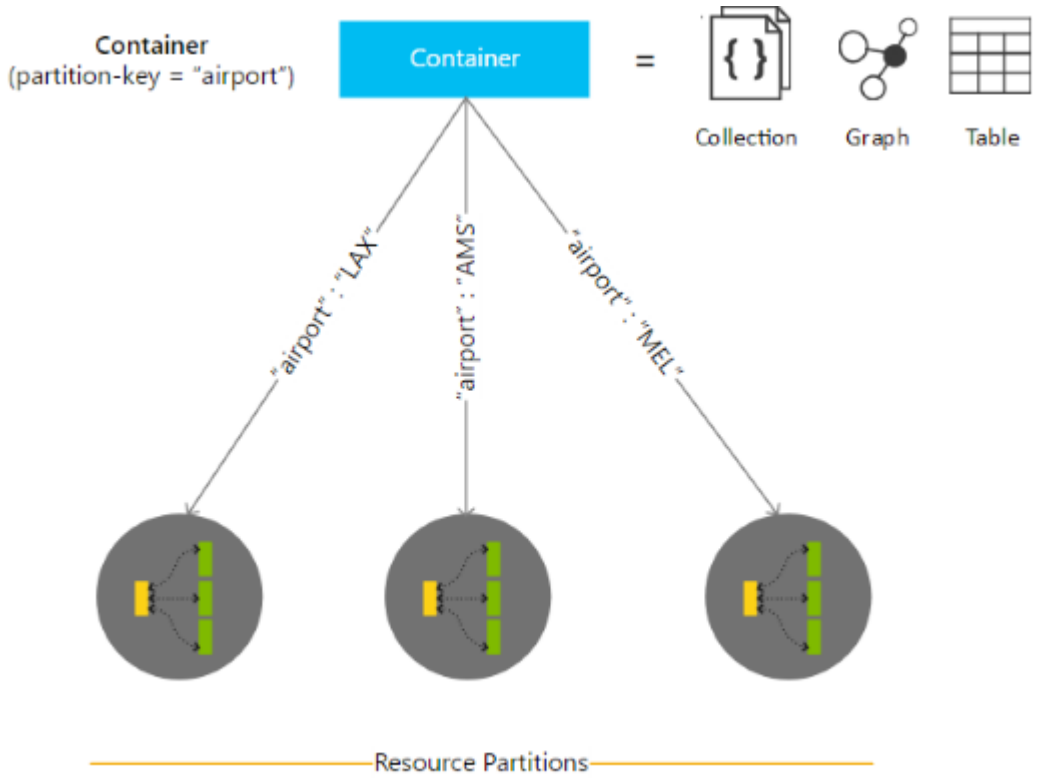

Different consistency levels

- 5 different types:
  - **Strong** -> **RU** ☹ 1 Azure region
  - Bounded Staleness ->reads behind writes
  - Session-> read your own writes guarantees
  - Consistent Prefix->reads never see out of order writes
  - **Eventual** -> the replicas eventually converge **RU** ☺
- Per database, collection or query

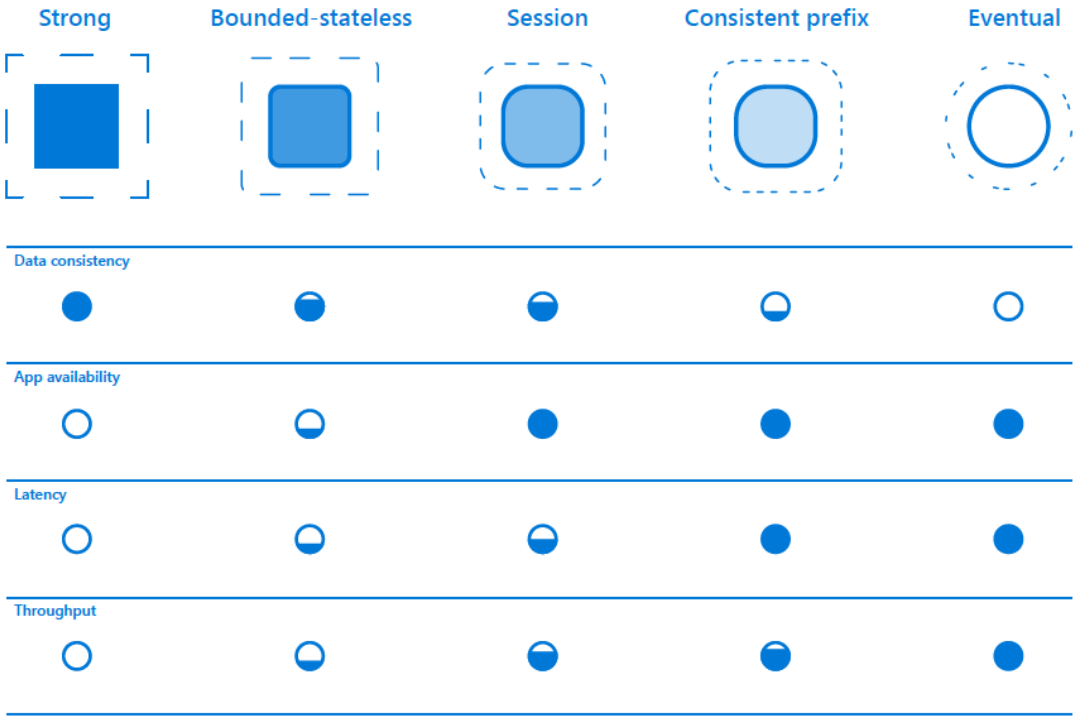

Customizable indexing

- All data is indexed automatically (without any assumptions or schema information)
- You can specify **custom indexing policies**
  - Include or exclude documents and paths to and from the index
  - Configure various index types
    - **Hash** -> equality
    - **Range** - > range queries
    - **Spatial** -> “distance” queries
  - Configure index update modes
    - **Consistent** -> updated synchronously as part of the document update
    - **Lazy** -> index is updated asynchronously
    - **None** -> Accessed only by their ID property

Using SQL with NoSQL

- It deals with hierarchical entities instead of rows and columns
- Schema-less data
- Relations across documents in a collection are captured by containment and not by primary key and foreign key relations

JSON

```
{
  "id": "WakefieldFamily",
  "parents": [
      { "familyName": "Wakefield", "givenName": "Robin" },
      { "familyName": "Miller", "givenName": "Ben" }
  ],
  "children": [
      {
        "familyName": "Merriam",
        "givenName": "Jesse",
        "gender": "female", "grade": 1,
        "pets": [
            { "givenName": "Goofy" },
            { "givenName": "Shadow" }
        ]
      },
      {
        "familyName": "Miller",
         "givenName": "Lisa",
         "gender": "female",
         "grade": 8 }
  ],
  "address": { "state": "NY", "county": "Manhattan", "city": "NY" },
  "creationDate": 1431620462,
  "isRegistered": false
}
```

```
SELECT *
FROM Families f
WHERE f.id = "AndersenFamily"
```

```
SELECT {"Name":f.id, "City":f.address.city} AS Family
FROM Families f
WHERE f.address.city = f.address.state
```

Time To Live Data

- Documents **automatically purged** after a period of time.
- You need to provide the period, in seconds, after the last modified timestamp of the document will be expired:
  - TTL per collection
  - TTL per document
- Expired document won't be accessible
- No impact on RU

Global distribution

- Low-latency access to data
- Business continuity and disaster recovery
- **1** Write region
- **N** Preferred read regions
- Manual\Automatic failover
- All Azure regions available
- Service\Client configuration

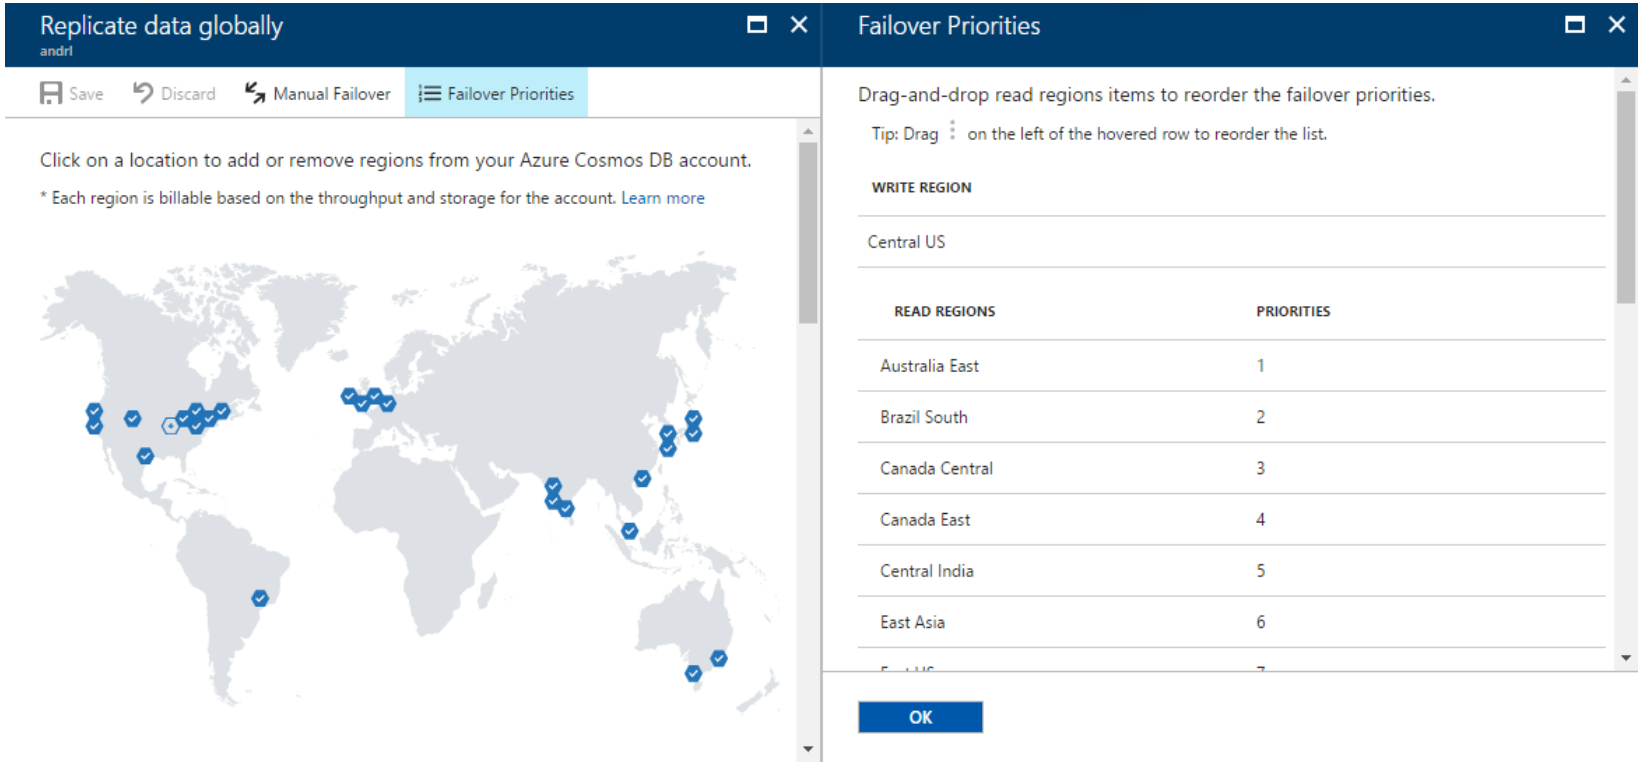

Development

Client side

- .NET
- Java
- JS
- Node.JS
- Pyton

```
var createDocumentStoredProc = {
    id: "createMyDocument",
    serverScript: function createMyDocument(documentToCreate) {
        var context = getContext();
        var collection = context.getCollection();

        var accepted = collection.createDocument(collection.getSelfLink(),
              documentToCreate,
              function (err, documentCreated) {
                  if (err) throw new Error('Error' + err.message);
                  context.getResponse().setBody(documentCreated.id)
              });
        if (!accepted) return;
    }
}
```

Server side

- Javascript
  - Stored Procedure -> transactional batch operations
  - Trigger -> Pre\Post
  - User Defined Function -> extend query language

Emulator

- Local environment that emulates the Azure Cosmos DB service
- For development purposes only
  - Single fixed account and a well-known master key
  - Not scalable
  - Not replicable
  - No consistency levels
  - No encryption
- Runs on https://localhost:8081/
- Command-line tool
- Docker support

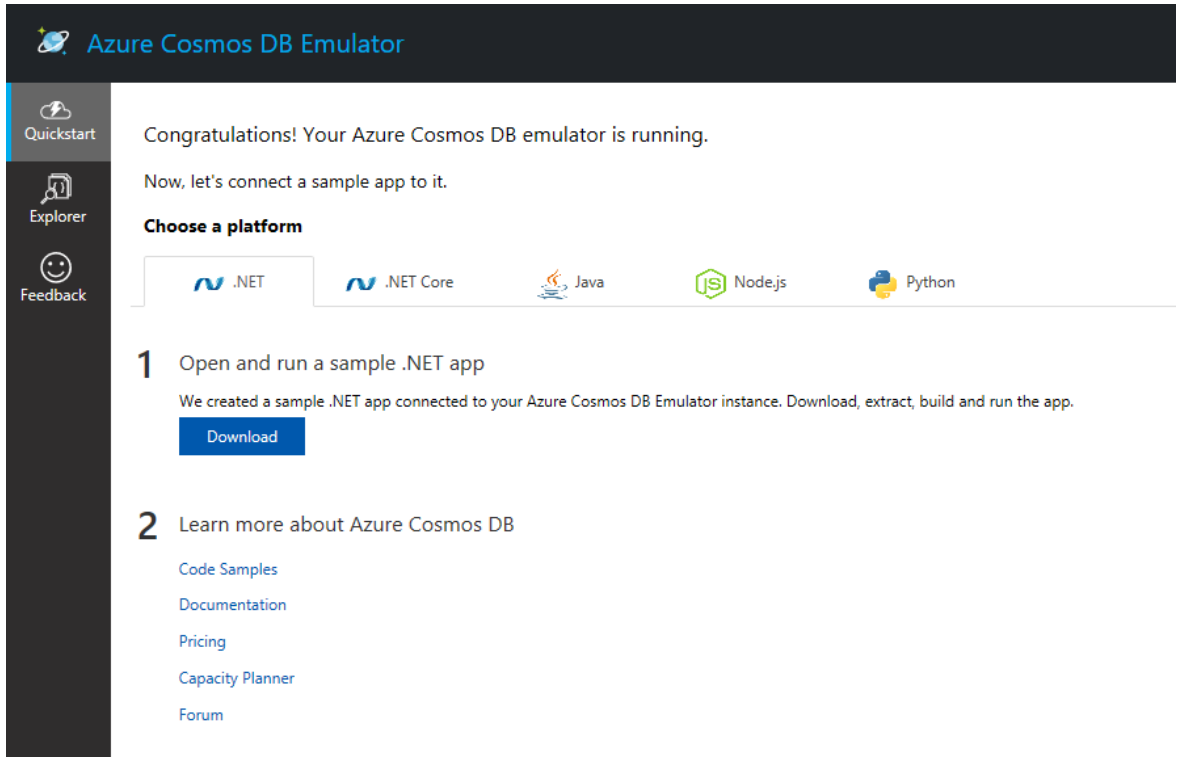

Tools

- **Azure Storage Explorer** https://azure.microsoft.com/en-us/features/storage-explorer/
- Data Migration Tool https://docs.microsoft.com/en-us/azure/cosmos-db/import-data
- DocumentDb Studio https://github.com/mingaliu/DocumentDBStudio
- Query Playground https://www.documentdb.com/sql/demo

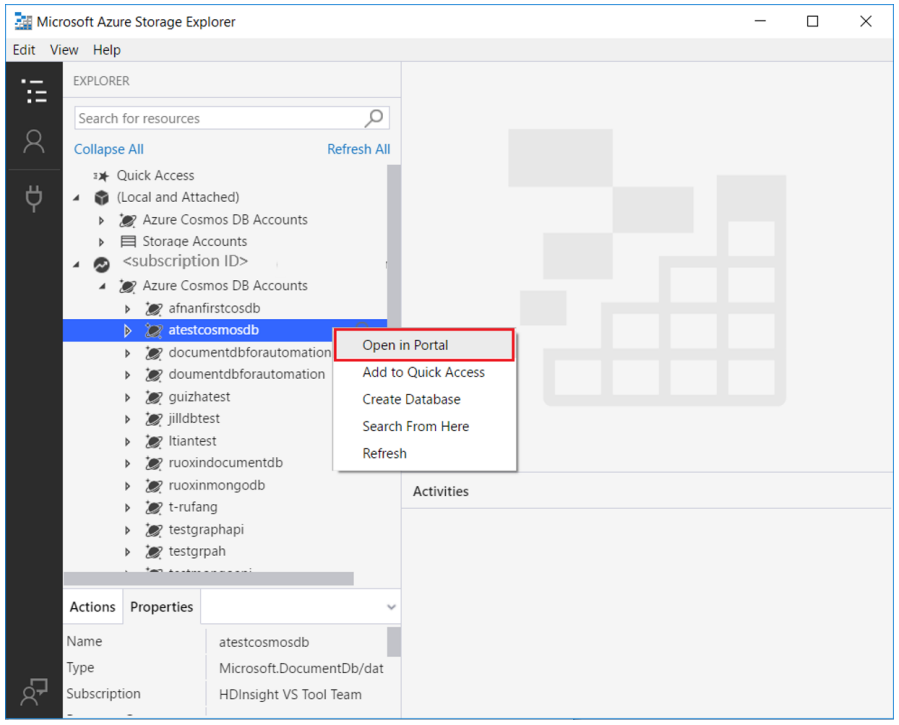

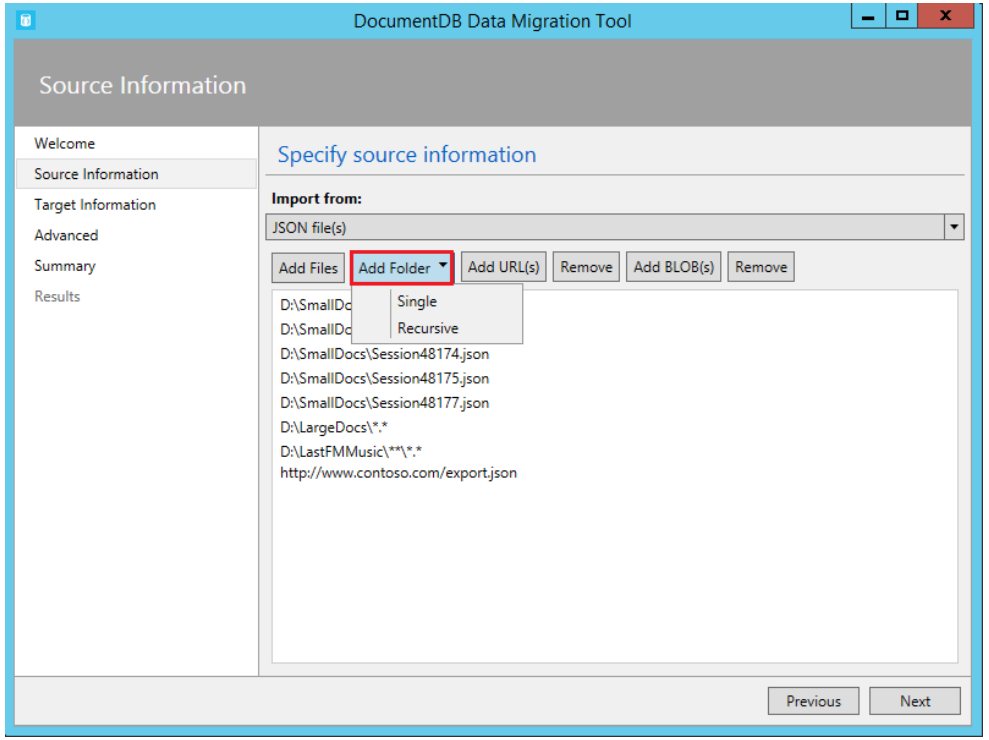

Backup & Restore

- Backups not affect the performance
- **Every four hours and latest 2 backups are stored ☹**
- **Contact Azure support ☹**

What's next

https://feedback.azure.com/forums/263030-azure-cosmos-db

- **Allow paging (skip/take)**
- **Partial update**
- GroupBy
- Distinct
- Batch operation (no SP)

Demo

https://github.com/martinobordin/CosmoGAB

Grazie

Domande?

 martinobordin

 @martinobordin

 martinobordin